

Assignment 1

Ariz Kazani

2024-05-31

Assignment 1

Name: Ariz Kazani

Student ID: 101311311

Notes

```
# TODO: add information about assignment and libraries used
#   - make sure to double check and double read each question
#   - format code ctr + shift + a
#   - test all code to make sure its still working
#   - use summaryfunction for dataframes where applicable
```

Solutions

Question 1. Function Creation and Vector Operations

(a) Create a vector named sales that contains the following sales figures for a week: 250, 310, 450, 500, 620, 715, and 840.

```
# vector that contains sales figures
sales <- c(250, 310, 450, 500, 620, 715, 840)
```

(b) Write a function named sales_summary that takes a vector as input and returns the sum and mean of the vector. Test your function using the sales vector.

```

sales_summary <- function(vector) {
  Sum <- sum(vector)
  Mean <- mean(vector)
  returnVar <- list(Sum = Sum, Mean = Mean)
  return(returnVar)
}

sales_summary(sales)

```

```

## $Sum
## [1] 3685
##
## $Mean
## [1] 526.4286

```

(c) Write a function named `adjust_sales` that takes a vector and a percentage as inputs, adjusts each entry in the vector by the given percentage, and returns the adjusted vector in descending order. Test your function with the sales vector and a 10% increase.

```

adjust_sales <- function(vector, percentage) {
  newVector <- vector * ((percentage + 100) / 100)
  newVector <- sort(newVector, decreasing = TRUE)
  return(newVector)
}

salesADJ <- adjust_sales(sales, 10)
salesADJ

```

```

## [1] 924.0 786.5 682.0 550.0 495.0 341.0 275.0

```

(d) Create another test for the `sales_summary` function with a random vector of 10 elements. Print the result to check if your function works correctly with different inputs.

```

randomVector = c(22, 3, 24, 536, 774678, 895676, 57635, 24344, 123, 534)

sales_summary(randomVector)

```

```

## $Sum
## [1] 1753575
##
## $Mean
## [1] 175357.5

```

(e) Similarly, test the `adjust_sales` function with a random vector of 10 elements and a random percentage between 5% and 20%. Print the adjusted vector to ensure your function works correctly.

```
adjust_sales(randomVector, 17)
```

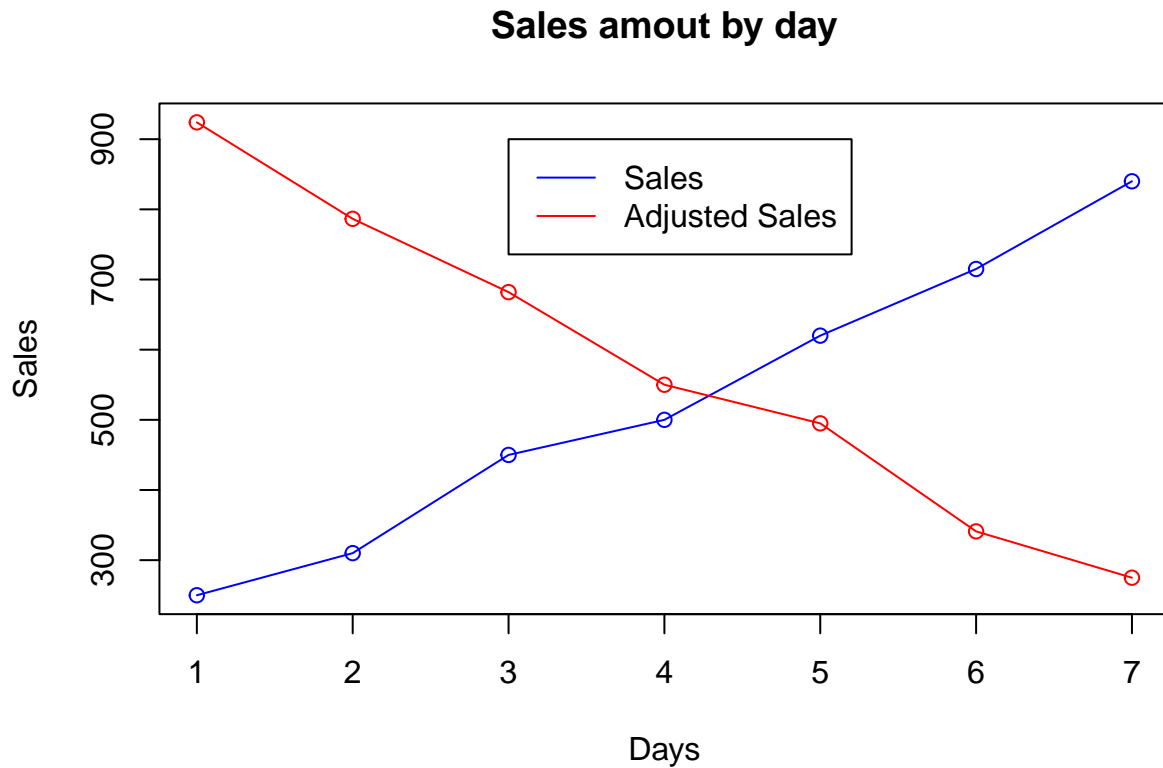
```
## [1] 1047940.92 906373.26 67432.95 28482.48 627.12 624.78
## [7] 143.91 28.08 25.74 3.51
```

(f) Plot the original sales vector and the adjusted sales vector (from Part 3) on the same graph using different colors. Label the axes and add a legend.

```
plot(
  sales,
  type = "o",
  col = "blue",
  ylim = range(c(sales, salesADJ)),
  main = "Sales amout by day",
  xlab = "Days",
  ylab = "Sales"
)

lines(salesADJ, type = "o", col = "red")

legend(
  3, 900,
  legend = c("Sales", "Adjusted Sales"),
  col = c("blue", "red"),
  lty = 1
)
```



Question 2. Dataframe Operations and Descriptive Statistics

(a) Create a dataframe named students with the following data:

- Name: "Alice", "Bob", "Charlie", "David", "Eva"
- Age: 23, 22, 24, 21, 23
- Score: 85, 92, 78, 88, 90

```
students = data.frame(
  Name = c("Alice", "Bob", "Charlie", "David", "Eva"),
  Age = c(23, 22, 24, 21, 23),
  Score = c(85, 92, 78, 88, 90)
)
```

students

```
##      Name Age Score
## 1  Alice  23   85
## 2   Bob  22   92
## 3 Charlie  24   78
## 4  David  21   88
## 5   Eva  23   90
```

(b) Add a new column to the students dataframe named Passed with a value of TRUE if the Score is 80 or above, and FALSE otherwise.

```
students$Passed <- students$Score >= 80
```

```
students
```

```
##      Name Age Score Passed
## 1  Alice  23    85    TRUE
## 2   Bob  22    92    TRUE
## 3 Charlie 24    78   FALSE
## 4  David 21    88    TRUE
## 5   Eva  23    90    TRUE
```

(c) Calculate the mean, median, and standard deviation of the Age and Score columns in the students dataframe.

```
ageData <- list(
  AgeMean = mean(students$Age),
  AgeMedian = median(students$Age),
  AgeSD = sd(students$Age)
)
```

```
scoreData <- list(
  ScoreMean = mean(students$Score),
  ScoreMedian = median(students$Score),
  ScoreSD = sd(students$Score)
)
```

```
ageData
```

```
## $AgeMean
## [1] 22.6
##
## $AgeMedian
## [1] 23
##
## $AgeSD
## [1] 1.140175
```

```
scoreData
```

```
## $ScoreMean
## [1] 86.6
##
## $ScoreMedian
## [1] 88
##
## $ScoreSD
## [1] 5.458938
```

(d) Identify the student(s) with the highest score and display their details.

```
maxScore = max(students$Score)

studentsWMS <- students[students$Score == maxScore, ]

studentsWMS
```

```
##   Name Age Score Passed
## 2   Bob  22    92   TRUE
```

(e) Filter the dataframe to show only the students who passed and save it as a new dataframe named `passed_students`.

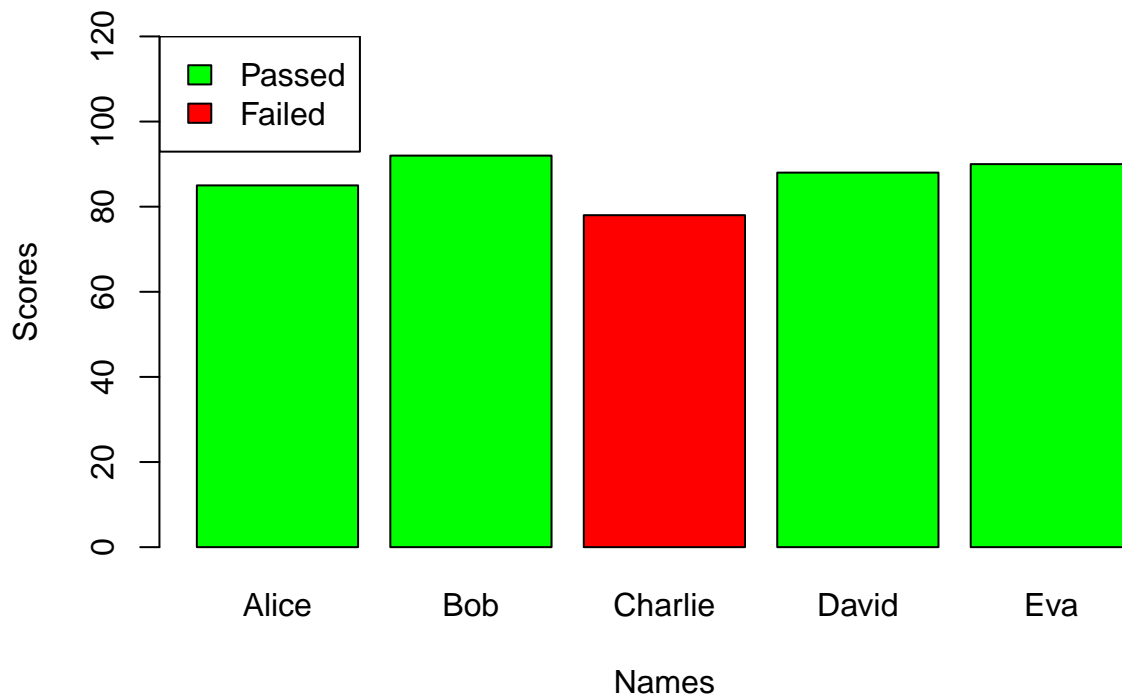
```
passed_students <- students[students$Passed, ]

passed_students
```

```
##   Name Age Score Passed
## 1 Alice  23    85   TRUE
## 2   Bob  22    92   TRUE
## 4 David  21    88   TRUE
## 5   Eva  23    90   TRUE
```

(f) Create a bar chart showing the scores of all students. Use different colors for those who passed and those who did not.

```
barplot(
  students$Score,
  names.arg = students$Name,
  ylim = c(0, 120),
  ylab = "Scores",
  xlab = "Names",
  col = ifelse(students$Passed, "Green", "Red"),
)
legend("topleft",
  legend = c("Passed", "Failed"),
  fill = c("Green", "Red"),)
```



(g) Write a short summary (3-5 sentences) interpreting the statistical results and the bar chart created in the previous steps.

From the presented data, we can conclude that most people passed, scoring well over the required 80%. We can also see that most people scored higher than the mean, with Alice and Charlie being the exceptions. The bar chart does not show any correlation between gender and the likelihood of passing. We can also see that most people are within one standard deviation from the mean, with Charlie and Bob being the exceptions.

3. Advanced Data Manipulation and Visualization

(a) Create a dataframe named `employees` with the following data:

- EmployeeID: 101, 102, 103, 104, 105
- Name: "John", "Jane", "Doe", "Smith", "Emily"
- Department: "Sales", "HR", "IT", "Finance", "Marketing"
- Salary: 60000, 65000, 70000, 72000, 68000
- Experience: 3, 7, 5, 10, 4

```
employees <- data.frame(
  EmployeeID = c(101, 102, 103, 104, 105),
  Name = c("John", "Jane", "Doe", "Smith", "Emily"),
  Department = c("Sales", "HR", "IT", "Finance", "Marketing"),
```

```
Salary = c(60000, 65000, 70000, 72000, 68000),
Experience = c(3, 7, 5, 10, 4)
)

employees
```

```
## EmployeeID Name Department Salary Experience
## 1      101  John      Sales  60000          3
## 2      102  Jane      HR    65000          7
## 3      103   Doe      IT    70000          5
## 4      104 Smith    Finance 72000         10
## 5      105 Emily Marketing 68000          4
```

(b) Calculate the mean and median salary for each department. Write a function named `department_summary` that returns a summary dataframe containing the department name, mean salary, and median salary.

```
department_summary <- function(dataFrame) {
  getMeanMedian <- function(df, name) {
    dfWithName <- df[df$Department == name, ]
    returnDF = data.frame(
      DepartmentName = name,
      MeanSalary = mean(dfWithName$Salary),
      MedianSalary = median(dfWithName$Salary)
    )
    return(summary(returnDF))
  }

  departments <- unique(dataFrame$Department)
  df = data.frame()
  for (dep in departments) {
    df <- rbind(df, getMeanMedian(dataFrame, dep))
  }
  return(df)
}

department_summary(employees)
```

```
## Var1      Var2      Freq
## 1 DepartmentName Length:1
## 2 DepartmentName Class :character
## 3 DepartmentName Mode  :character
## 4 DepartmentName      <NA>
## 5 DepartmentName      <NA>
## 6 DepartmentName      <NA>
## 7      MeanSalary Min.   :60000
## 8      MeanSalary 1st Qu.:60000
## 9      MeanSalary Median :60000
## 10     MeanSalary Mean   :60000
## 11     MeanSalary 3rd Qu.:60000
## 12     MeanSalary Max.   :60000
```



```

## 13      MedianSalary      Min.   :60000
## 14      MedianSalary      1st Qu.:60000
## 15      MedianSalary      Median :60000
## 16      MedianSalary      Mean   :60000
## 17      MedianSalary      3rd Qu.:60000
## 18      MedianSalary      Max.   :60000
## 19      DepartmentName Length:1
## 20      DepartmentName Class :character
## 21      DepartmentName Mode  :character
## 22      DepartmentName      <NA>
## 23      DepartmentName      <NA>
## 24      DepartmentName      <NA>
## 25      MeanSalary      Min.   :65000
## 26      MeanSalary      1st Qu.:65000
## 27      MeanSalary      Median :65000
## 28      MeanSalary      Mean   :65000
## 29      MeanSalary      3rd Qu.:65000
## 30      MeanSalary      Max.   :65000
## 31      MedianSalary      Min.   :65000
## 32      MedianSalary      1st Qu.:65000
## 33      MedianSalary      Median :65000
## 34      MedianSalary      Mean   :65000
## 35      MedianSalary      3rd Qu.:65000
## 36      MedianSalary      Max.   :65000
## 37      DepartmentName Length:1
## 38      DepartmentName Class :character
## 39      DepartmentName Mode  :character
## 40      DepartmentName      <NA>
## 41      DepartmentName      <NA>
## 42      DepartmentName      <NA>
## 43      MeanSalary      Min.   :70000
## 44      MeanSalary      1st Qu.:70000
## 45      MeanSalary      Median :70000
## 46      MeanSalary      Mean   :70000
## 47      MeanSalary      3rd Qu.:70000
## 48      MeanSalary      Max.   :70000
## 49      MedianSalary      Min.   :70000
## 50      MedianSalary      1st Qu.:70000
## 51      MedianSalary      Median :70000
## 52      MedianSalary      Mean   :70000
## 53      MedianSalary      3rd Qu.:70000
## 54      MedianSalary      Max.   :70000
## 55      DepartmentName Length:1
## 56      DepartmentName Class :character
## 57      DepartmentName Mode  :character
## 58      DepartmentName      <NA>
## 59      DepartmentName      <NA>
## 60      DepartmentName      <NA>
## 61      MeanSalary      Min.   :72000
## 62      MeanSalary      1st Qu.:72000
## 63      MeanSalary      Median :72000
## 64      MeanSalary      Mean   :72000
## 65      MeanSalary      3rd Qu.:72000
## 66      MeanSalary      Max.   :72000

```

```
## 67      MedianSalary   Min.   :72000
## 68      MedianSalary   1st Qu.:72000
## 69      MedianSalary   Median :72000
## 70      MedianSalary   Mean    :72000
## 71      MedianSalary   3rd Qu.:72000
## 72      MedianSalary   Max.    :72000
## 73      DepartmentName Length:1
## 74      DepartmentName Class   :character
## 75      DepartmentName Mode    :character
## 76      DepartmentName      <NA>
## 77      DepartmentName      <NA>
## 78      DepartmentName      <NA>
## 79      MeanSalary     Min.    :68000
## 80      MeanSalary     1st Qu.:68000
## 81      MeanSalary     Median   :68000
## 82      MeanSalary     Mean     :68000
## 83      MeanSalary     3rd Qu.:68000
## 84      MeanSalary     Max.     :68000
## 85      MedianSalary   Min.    :68000
## 86      MedianSalary   1st Qu.:68000
## 87      MedianSalary   Median   :68000
## 88      MedianSalary   Mean     :68000
## 89      MedianSalary   3rd Qu.:68000
## 90      MedianSalary   Max.     :68000
```

(c) Identify and display details of the employee with the highest salary in each department. Write a function named `top_earner` to achieve this.

```
top_earner <- function(dataFrame) {
  getMax <- function(df, dep) {
    peopleByDep <- df[df$Department == dep, ]
    maxSal <- max(peopleByDep$Salary)
    topEarner <- peopleByDep[peopleByDep$Salary == maxSal, ]
    return(topEarner)
  }
  departments <- unique(dataFrame$Department)
  df = data.frame()
  for (dep in departments) {
    df <- rbind(df, getMax(dataFrame, dep))
  }
  return(df)
}

top_earner(employees)
```

```
##   EmployeeID  Name Department Salary Experience
## 1         101  John      Sales  60000          3
## 2         102  Jane       HR   65000          7
## 3         103   Doe       IT   70000          5
## 4         104 Smith  Finance  72000         10
## 5         105 Emily Marketing 68000          4
```

(d) Add a new column to the employees dataframe named AdjustedSalary, which is the Salary adjusted for experience (increase by 2% for each year of experience).

```
employees$AdjustedSalary <- employees$Salary *  
  ((100 + employees$Experience * 2) / 100)  
  
employees
```

	EmployeeID	Name	Department	Salary	Experience	AdjustedSalary
## 1	101	John	Sales	60000	3	63600
## 2	102	Jane	HR	65000	7	74100
## 3	103	Doe	IT	70000	5	77000
## 4	104	Smith	Finance	72000	10	86400
## 5	105	Emily	Marketing	68000	4	73440

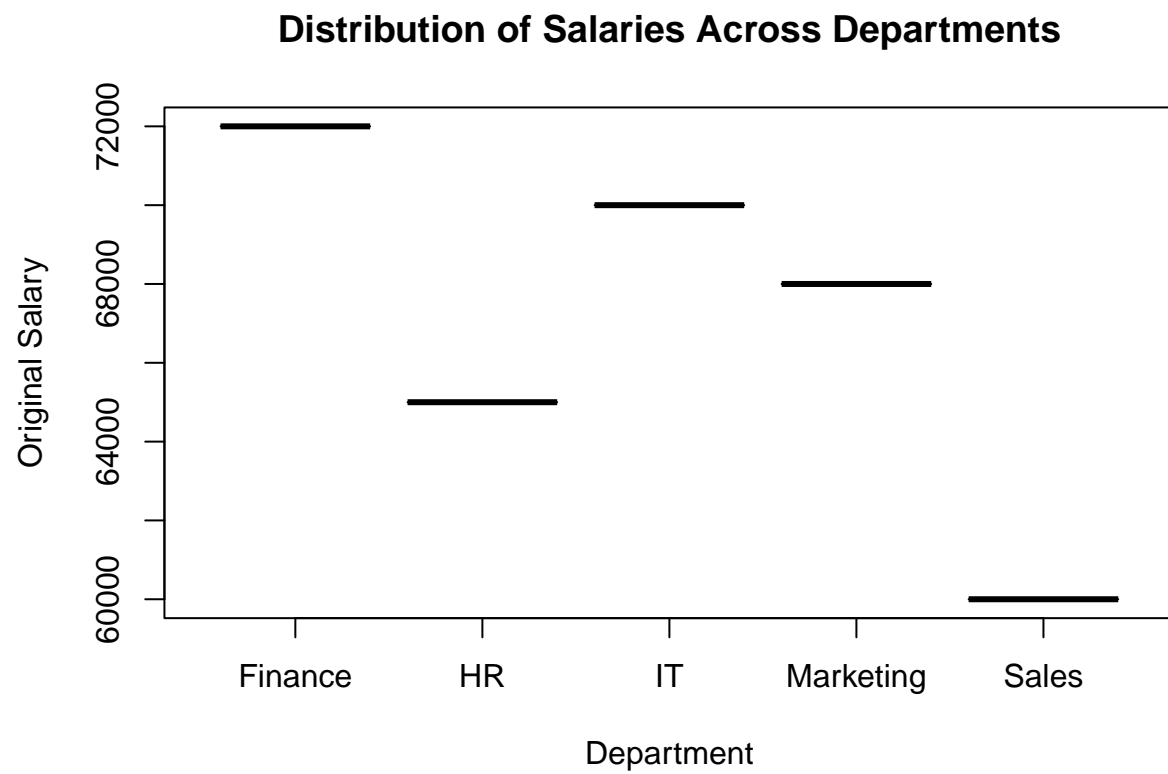
(e) Filter the dataframe to show only employees with an adjusted salary above 70,000 and save it as a new dataframe named high_earners.

```
high_earners <- employees[employees$AdjustedSalary > 70000, ]  
  
high_earners
```

	EmployeeID	Name	Department	Salary	Experience	AdjustedSalary
## 2	102	Jane	HR	65000	7	74100
## 3	103	Doe	IT	70000	5	77000
## 4	104	Smith	Finance	72000	10	86400
## 5	105	Emily	Marketing	68000	4	73440

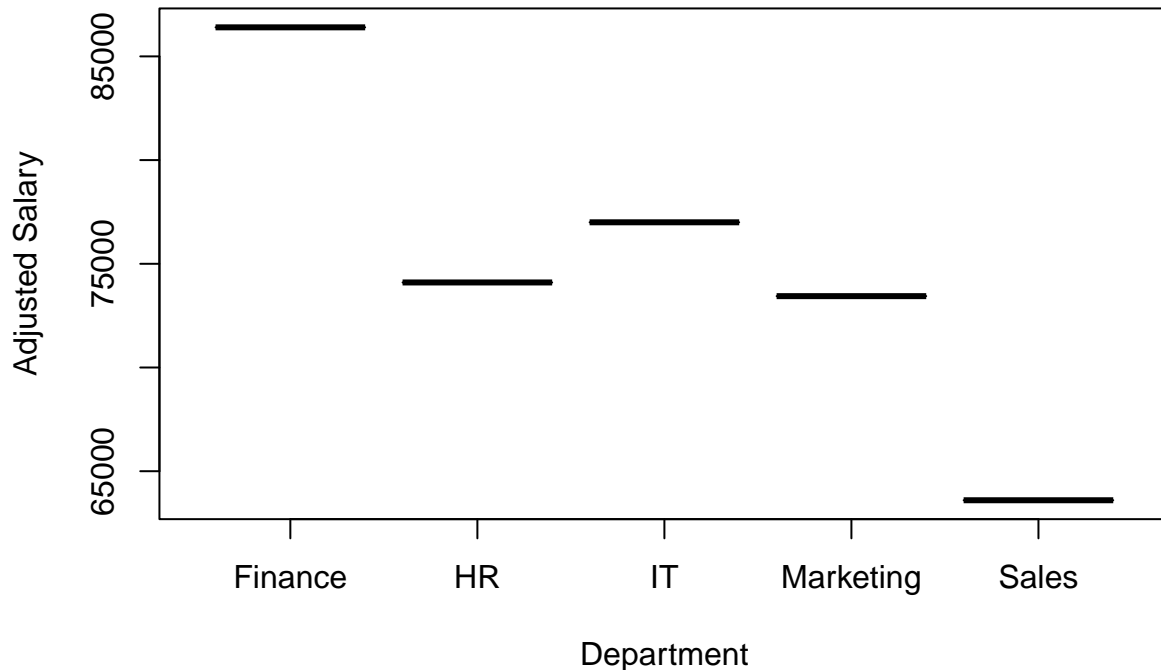
(f) Create a boxplot to compare the distribution of original salaries and adjusted salaries across different departments. Add appropriate labels and a title.

```
boxplot(  
  employees$Salary ~ employees$Department,  
  main = "Distribution of Salaries Across Departments",  
  xlab = "Department",  
  ylab = "Original Salary",  
)
```



```
boxplot(  
  employees$AdjustedSalary ~ employees$Department,  
  main = "Distribution of Adjusted Salaries Across Departments",  
  xlab = "Department",  
  ylab = "Adjusted Salary",  
)
```

Distribution of Adjusted Salaries Across Departments



(g) Write a short analysis (4-6 sentences) interpreting the results from the summary statistics, top earners, and the boxplot.

From the statistics we can see that the mean and median are exactly the same because there is only one person in each department. From the top-earners we can see that most people are earning above 70000 with the exception being sales. This is likely the case as sales has the lowest starting salary and the person working sales has been there the least amount of time. From the box plots we can see that the mean of the salary between all departments has increased, with the biggest jump happening from finance. We can also see, with the adjustment HR's mean salary goes from 4th to 3rd.

4. Exploring Dataframes with Multiple Operations

(a) Create a dataframe named products with the following data:

- ProductID: 201, 202, 203, 204, 205
- ProductName: "Laptop", "Smartphone", "Tablet", "Headphones", "Smartwatch"
- Category: "Electronics", "Electronics", "Electronics", "Accessories", "Electronics"
- Price: 1200, 800, 600, 200, 350
- QuantitySold: 150, 200, 300, 400, 250

```
products <- data.frame(
  ProductID = c(201, 202, 203, 204, 205),
  ProductName = c("Laptop", "Smartphone", "Tablet", "Headphones", "Smartwatch"),
  Category = c(
    "Electronics",
    "Electronics",
    "Electronics",
    "Accessories",
    "Electronics"
  ),
  Price = c(1200, 800, 600, 200, 350),
  QuantitySold = c(150, 200, 300, 400, 250)
)
```

(b) Calculate the total revenue for each product (Price * QuantitySold). Write a function named `calculate_revenue` that adds a new column `Revenue` to the `products` dataframe.

```
calculate_revenue <- function(dataFrame) {
  dataFrame$Revenue <- dataFrame$Price * dataFrame$QuantitySold
  return(dataFrame)
}

products <- calculate_revenue(products)

products
```

##	ProductID	ProductName	Category	Price	QuantitySold	Revenue
## 1	201	Laptop	Electronics	1200	150	180000
## 2	202	Smartphone	Electronics	800	200	160000
## 3	203	Tablet	Electronics	600	300	180000
## 4	204	Headphones	Accessories	200	400	80000
## 5	205	Smartwatch	Electronics	350	250	87500

(c) Identify the product with the highest revenue and display its details.

```
maxRevenue <- max(products$Revenue)
products[products$Revenue == maxRevenue, ]
```

##	ProductID	ProductName	Category	Price	QuantitySold	Revenue
## 1	201	Laptop	Electronics	1200	150	180000
## 3	203	Tablet	Electronics	600	300	180000

(d) Group the products by `Category` and calculate the total revenue for each category. Write a function named `category_revenue` that returns a summary dataframe with `Category` and `TotalRevenue`.

```

category_revenue <- function(dataFrame) {
  dataByProduct <- function(dataFrame, category) {
    toteRev <- sum(dataFrame[dataFrame$Category == category, 'Revenue'])
    returnDf <- data.frame(Category = category, TotalRevenue = toteRev)
    return(returnDf)
  }
  categories <- unique(dataFrame$Category)
  df <- data.frame()
  for (cat in categories) {
    df <- rbind(df, dataByProduct(dataFrame, cat))
  }
  return(df)
}

category_revenue(products)

```

```

##      Category TotalRevenue
## 1 Electronics      607500
## 2 Accessories      80000

```

(e) Create a bar chart to display the total revenue for each product, and use different colors for each category.

```


# Create a list with categories as names and colors as values
categories <- unique(products$Category)
numCat <- length(categories)
colours <- rainbow(numCat)

colourList <- list()
for (i in 1:numCat) {
  colourList[[categories[i]]] <- colours[i]
}

barplot(
  products$Revenue,
  names.arg = products$Name,
  ylab = "Scores",
  xlab = "Names",
  col = sapply(products$Category, function(col) colourList[[col]])
)

legend("topright",
  legend = categories,
  fill = colours)

```



(f) Generate a scatter plot of Price versus QuantitySold with different colors for each category. Add a trend line to the plot.

```
# TODO: fill this out
```

(g) Write a detailed report (5-7 sentences) analyzing the

5. Debugging Subsetting and Indexing Issues

Explain the issues with the code and provide the correct working code. Output the code to show that you have it corrected.

(a)

```
students <- data.frame(
  Name = c("Alice", "Bob", "Charlie", "David", "Eva"),
  Age = c(23, 22, 24, 21, 23),
  Score = c(85, 92, 78, 88, 90))
```



```
# Extracting ages of students who scored above 80
high_scorers_ages <- students[students$Score > 80][, "Age"]
print(high_scorers_ages)
```

Explanation: `students[students$Score > 80]` is not correctly specifying rows and columns, to fix this we need to add a comma after specifying the columns like this `students[students$Score > 80,]`. Although it is not syntactically wrong we can also combine `students[students$Score > 80,]` and `[, "Age"]` to look more clean `students[students$Score > 80, "Age"]`. **Corrected Code:**

```
students <- data.frame(
  Name = c("Alice", "Bob", "Charlie", "David", "Eva"),
  Age = c(23, 22, 24, 21, 23),
  Score = c(85, 92, 78, 88, 90))
# Extracting ages of students who scored above 80
high_scorers_ages <- students[students$Score > 80, "Age"]
print(high_scorers_ages)
```

```
## [1] 23 22 21 23
```

(b)

```
employee_list <- list(
  Name = "John",
  Age = 30,
  Department = "HR",
  Salary = 50000
)
# Accessing the salary of the employee
salary <- employee_list["Salaries"]
print(salary)
```

Explanation: The salary of the employee is not correctly being accessed. It was assigned as `Salary` but we are attempting to get `"Salaries"`. Instead we need to get `"Salary"`. **Corrected Code:**

```
employee_list <- list(
  Name = "John",
  Age = 30,
  Department = "HR",
  Salary = 50000
)
# Accessing the salary of the employee
salary <- employee_list["Salary"]
print(salary)
```

```
## $Salary
## [1] 50000
```

(c)

```

sales_data <- array(1:27, dim = c(3, 3, 3))
# Extracting the value in the second row, second column of the first matrix
value <- sales_data[3, 3, 0]
print(value)

```

Explanation: The first issue is that we aren't correctly selecting the row and column if we want the second row and second column we want to get the data from [2, 2,]. the second issue is that we are not correctly accessing the first matrix, since R is 1 indexed and not 0, to get the first matrix we need to use 1 to access it and not 0: sales_data[2, 2, 1]. **Corrected Code:**

```

sales_data <- array(1:27, dim = c(3, 3, 3))
# Extracting the value in the second row, second column of the first matrix
value <- sales_data[2, 2, 1]
print(value)

```

```
## [1] 5
```

(d)

```

products <- data.frame(
  ProductID = c(201, 202, 203, 204, 205),
  ProductName = c("Laptop", "Smartphone", "Tablet", "Headphones", "Smartwatch"),
  Category = c("Electronics", "Electronics", "Electronics", "Accessories", "Electronics"),
  Price = c(1200, 800, 600, 200, 350),
  QuantitySold = c(150, 200, 300, 400, 250)
)
# Extracting products with a price above 500
expensive_products <- products[products$Price >= "500", ]
print(expensive_products)

```

Explanation: The first issue is that we are trying to compare string value with a numeric value. When this happens the numeric value gets converted into a string and gets compared lexicographically. This will not always give the desired result, for example 1200 gets converted into "1200" and then gets compared with "500". Since "1" is smaller than "5" it will not add it. Another issue is that we want to get prices above 500 but we are also including 500 which is not what we want. **Corrected Code:**

```

products <- data.frame(
  ProductID = c(201, 202, 203, 204, 205),
  ProductName = c("Laptop", "Smartphone", "Tablet", "Headphones", "Smartwatch"),
  Category = c("Electronics", "Electronics", "Electronics", "Accessories", "Electronics"),
  Price = c(1200, 800, 600, 200, 350),
  QuantitySold = c(150, 200, 300, 400, 250)
)
# Extracting products with a price above 500
expensive_products <- products[products$Price > 500, ]
print(expensive_products)

```

```

##   ProductID ProductName   Category Price QuantitySold
## 1       201     Laptop Electronics 1200           150
## 2       202  Smartphone Electronics   800           200
## 3       203      Tablet Electronics   600           300

```

6. Analysis of the “trees” Dataset

This dataset has three variables (Girth, Height, Volume) on 31 felled black cherry trees.

(a)

- Load the “trees” dataset and check the structure with `str()`.
- Use `apply()` to return the mean values for the three variables (Girth, Height, Volume) and output these values.
- Determine the number of trees with Volume greater than the mean Volume.

```
data(trees)
str(trees)
```

```
## 'data.frame':  31 obs. of  3 variables:
## $ Girth : num  8.3 8.6 8.8 10.5 10.7 10.8 11 11 11.1 11.2 ...
## $ Height: num  70 65 63 72 81 83 66 75 80 75 ...
## $ Volume: num  10.3 10.3 10.2 16.4 18.8 19.7 15.6 18.2 22.6 19.9 ...
```

```
meanValues <- apply(trees, 2, mean)
meanValues
```

```
##      Girth      Height      Volume
## 13.24839 76.00000 30.17097
```

```
meanVolume <- meanValues["Volume"]
treesBiggerMean <- sum(trees$Volume > meanVolume)
treesBiggerMean
```

```
## [1] 12
```

(b)

- Convert each Girth (diameter) to a radius r .
- Calculate the cross-sectional area of each tree using $\pi \times r^2$.
- Calculate and output the interquartile range (IQR) of the areas.

```
trees$Girth <- trees$Girth / 2

crossSecArea <- (pi * trees$Girth ** 2)

IQRCSA <- IQR(crossSecArea)
IQRCSA
```

```
## [1] 87.1949
```

(c)

- Create a histogram of the areas calculated in part (b).
- Title and label the axes.

```
# TODO: sdfdmlsd
```

(d)

- Identify the tree with the largest area.
- Output its row number and the three measurements (Girth, Height, Volume) on one line

```
# TODO: sdfdmlsd
```

7. Comprehensive Data Analysis and Function Creation

(a)

- Load the mtcars dataset.
- Filter the dataset to include only cars with 6 or more cylinders and horsepower greater than 150. Save this filtered dataset as filtered_cars.

```
data(mtcars)

filtered_cars <- mtcars[mtcars$cyl >= 6 & mtcars$hp > 150, ]

filtered_cars
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
## Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
## Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
## Merc 450SE	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3
## Merc 450SL	17.3	8	275.8	180	3.07	3.730	17.60	0	0	3	3
## Merc 450SLC	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3	3
## Cadillac Fleetwood	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4
## Lincoln Continental	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4
## Chrysler Imperial	14.7	8	440.0	230	3.23	5.345	17.42	0	0	3	4
## Camaro Z28	13.3	8	350.0	245	3.73	3.840	15.41	0	0	3	4
## Pontiac Firebird	19.2	8	400.0	175	3.08	3.845	17.05	0	0	3	2
## Ford Pantera L	15.8	8	351.0	264	4.22	3.170	14.50	0	1	5	4
## Ferrari Dino	19.7	6	145.0	175	3.62	2.770	15.50	0	1	5	6
## Maserati Bora	15.0	8	301.0	335	3.54	3.570	14.60	0	1	5	8

(b)

- Create a function named efficiency_score that calculates an efficiency score for each car based on the formula: $EfficiencyScore = \frac{mpg}{(hp \times wt)}$

- Apply this function to the `filtered_cars` dataset and add the resulting scores as a new column named `Efficiency`.

```
efficiency_score <- function(dataFrame) {
  efs <- dataFrame$mpg / (dataFrame$hp * dataFrame$wt)
  return(efs)
}

filtered_cars$Efficiency <- efficiency_score(filtered_cars)
filtered_cars
```

```
##           mpg cyl  disp  hp drat   wt  qsec vs am gear carb
## Hornet Sportabout  18.7   8 360.0 175 3.15 3.440 17.02  0  0   3   2
## Duster 360        14.3   8 360.0 245 3.21 3.570 15.84  0  0   3   4
## Merc 450SE        16.4   8 275.8 180 3.07 4.070 17.40  0  0   3   3
## Merc 450SL        17.3   8 275.8 180 3.07 3.730 17.60  0  0   3   3
## Merc 450SLC       15.2   8 275.8 180 3.07 3.780 18.00  0  0   3   3
## Cadillac Fleetwood 10.4   8 472.0 205 2.93 5.250 17.98  0  0   3   4
## Lincoln Continental 10.4   8 460.0 215 3.00 5.424 17.82  0  0   3   4
## Chrysler Imperial 14.7   8 440.0 230 3.23 5.345 17.42  0  0   3   4
## Camaro Z28        13.3   8 350.0 245 3.73 3.840 15.41  0  0   3   4
## Pontiac Firebird   19.2   8 400.0 175 3.08 3.845 17.05  0  0   3   2
## Ford Pantera L     15.8   8 351.0 264 4.22 3.170 14.50  0  1   5   4
## Ferrari Dino       19.7   6 145.0 175 3.62 2.770 15.50  0  1   5   6
## Maserati Bora      15.0   8 301.0 335 3.54 3.570 14.60  0  1   5   8
##           Efficiency
## Hornet Sportabout  0.031063123
## Duster 360        0.016349397
## Merc 450SE        0.022386022
## Merc 450SL        0.025767054
## Merc 450SLC       0.022339800
## Cadillac Fleetwood 0.009663182
## Lincoln Continental 0.008918159
## Chrysler Imperial  0.011957539
## Camaro Z28        0.014136905
## Pontiac Firebird   0.028534275
## Ford Pantera L     0.018879648
## Ferrari Dino       0.040639505
## Maserati Bora      0.012542330
```

(c)

- Identify rows where the Efficiency score is less than the 1st percentile or greater than the 99th percentile of all Efficiency scores.
- Replace these outlier values with the mean Efficiency score of the remaining cars.

```
percential <- quantile(filtered_cars$Efficiency, probs = c(0.01, 0.99))

outLiers <- which(filtered_cars$Efficiency < percential[1] | filtered_cars$Efficiency > percential[2])

meanRC <- mean(filtered_cars$Efficiency[-outLiers])
```

```
filtered_cars$Efficiency[outLiers] <- meanRC
filtered_cars
```

```
##          mpg cyl  disp  hp drat   wt  qsec vs am gear carb
## Hornet Sportabout  18.7   8 360.0 175 3.15 3.440 17.02  0  0   3   2
## Duster 360        14.3   8 360.0 245 3.21 3.570 15.84  0  0   3   4
## Merc 450SE        16.4   8 275.8 180 3.07 4.070 17.40  0  0   3   3
## Merc 450SL        17.3   8 275.8 180 3.07 3.730 17.60  0  0   3   3
## Merc 450SLC       15.2   8 275.8 180 3.07 3.780 18.00  0  0   3   3
## Cadillac Fleetwood 10.4   8 472.0 205 2.93 5.250 17.98  0  0   3   4
## Lincoln Continental 10.4   8 460.0 215 3.00 5.424 17.82  0  0   3   4
## Chrysler Imperial 14.7   8 440.0 230 3.23 5.345 17.42  0  0   3   4
## Camaro Z28        13.3   8 350.0 245 3.73 3.840 15.41  0  0   3   4
## Pontiac Firebird   19.2   8 400.0 175 3.08 3.845 17.05  0  0   3   2
## Ford Pantera L     15.8   8 351.0 264 4.22 3.170 14.50  0  1   5   4
## Ferrari Dino       19.7   6 145.0 175 3.62 2.770 15.50  0  1   5   6
## Maserati Bora      15.0   8 301.0 335 3.54 3.570 14.60  0  1   5   8
##          Efficiency
## Hornet Sportabout  0.031063123
## Duster 360        0.016349397
## Merc 450SE        0.022386022
## Merc 450SL        0.025767054
## Merc 450SLC       0.022339800
## Cadillac Fleetwood 0.009663182
## Lincoln Continental 0.019419934
## Chrysler Imperial  0.011957539
## Camaro Z28        0.014136905
## Pontiac Firebird   0.028534275
## Ford Pantera L     0.018879648
## Ferrari Dino       0.019419934
## Maserati Bora      0.012542330
```

(d)

- Create a scatter plot of hp versus Efficiency, with points colored by the number of cylinders (cyl).
- Add a trend line to the scatter plot.
- Write a detailed analysis (6-8 sentences) interpreting the relationship between horsepower and efficiency, considering the number of cylinders.

```
# TODO: sdfdmLsd
```